

## Questions for preparation of T1 (OS 2023)

--

List all the blocks/data-structures on the ext2 partition (e.g. superblock, group descriptor, block bitmap, directory data block, etc) with block size 4KB, that MAY be updated (means changed) when a file of size 8 KB is created.

**Superblock:** The superblock contains information about the file system, such as the block size, inode count, and block count. When the block size changes, the superblock must be updated to reflect the new block size.

**Block Group Descriptor Table:** The block group descriptor table contains information about each block group in the file system, including the number of blocks and inodes in each block group. When the block size changes, the block group descriptor table must be updated to reflect the new block size.

**Inode Table:** The inode table contains information about each file and directory in the file system. When the block size changes, the inode table must be updated to reflect the new block size.

**Data Blocks:** The data blocks contain the actual file and directory data. When the block size changes, the data blocks must be reorganized and possibly moved to different locations on the partition to align with the new block size.

**Block Bitmap:** The block bitmap tracks which blocks are in use and which blocks are free. When the block size changes, the block bitmap must be updated to reflect the new block size.

**Inode Bitmap:** The inode bitmap tracks which inodes are in use and which inodes are free. When the block size changes, the inode bitmap must be updated to reflect the new block size.

List any 3 problems of each of the block allocation schemes (continuous, linked, indexed)

The continuous block allocation scheme is a file allocation scheme in which files are stored on a disk in contiguous blocks. While this scheme has some advantages, it also has several problems, including:

**External Fragmentation:** As files are added, deleted, and modified, the free space on the disk becomes fragmented into smaller and smaller pieces, making it difficult to allocate contiguous blocks for new files. This can lead to wasted space and reduced performance.

**Limited File Size:** With continuous block allocation, the maximum file size is limited by the size of the largest contiguous block of free space on the disk. If the disk becomes heavily fragmented, it may not be possible to allocate a contiguous block large enough to hold a large file.

**Slow File Allocation:** With continuous block allocation, finding a large enough contiguous block of free space to allocate to a new file can be time-consuming, especially on a heavily fragmented disk. This can lead to slower file allocation times and reduced performance.

**Difficult to Extend Files:** With continuous block allocation, extending the size of a file can be difficult, especially if there is no contiguous free space immediately following the file's current location. This may require moving the entire file to a new location on the disk, which can be time-consuming and may require significant disk I/O operations.

**Disk Defragmentation Required:** To overcome the issue of external fragmentation, the disk needs to be defragmented frequently. This process can be time-consuming and may require a significant amount of disk I/O operations, which can result in reduced system performance.

**Low Disk Utilization:** Continuous block allocation can lead to low disk utilization, as it may not be possible to allocate smaller portions of the disk to files. This can result in a significant amount of wasted space on the disk.

The linked block allocation scheme is a file allocation scheme in which files are stored on a disk using linked blocks, where each block contains a pointer to the next block in the file. While this scheme has some advantages, it also has several problems, including:

**High Overhead:** With linked block allocation, each block in the file requires an additional pointer to the next block, which can result in a significant amount of overhead. This can increase the size of the file and require more disk I/O operations, which can lead to reduced performance.

**Poor Performance for Random Access:** With linked block allocation, accessing a specific block in the file requires following a chain of pointers from the beginning of the file to the desired block. This can be slow and inefficient, especially for random access to non-contiguous blocks in the file.

**Disk Fragmentation:** With linked block allocation, small files may be scattered across the disk, leading to fragmentation. This can result in wasted space and reduced performance, as the disk needs to seek between non-contiguous blocks to read and write files.

**Limited Concurrency:** With linked block allocation, concurrent access to the same file can be difficult, especially if multiple processes are trying to modify the same block in the file. This can result in data inconsistencies and reduced performance.

**Security Vulnerabilities:** Linked block allocation can also be vulnerable to security attacks, such as buffer overflow attacks, in which an attacker can modify the pointers in a file to point to malicious code. This can allow the attacker to execute arbitrary code on the system or gain unauthorized access to sensitive data.

**Disk Space Overhead:** Linked block allocation also requires additional disk space to store the pointers between blocks, which can reduce the overall disk space utilization and waste disk space.

The linked block allocation scheme is a file allocation scheme in which files are stored on a disk using linked blocks, where each block contains a pointer to the next block in the file. While this scheme has some advantages, it also has several problems, including:

**High Overhead:** With linked block allocation, each block in the file requires an additional pointer to the next block, which can result in a significant amount of overhead. This can increase the size of the file and require more disk I/O operations, which can lead to reduced performance.

**Poor Performance for Random Access:** With linked block allocation, accessing a specific block in the file requires following a chain of pointers from the beginning of the file to the desired block. This can be slow and inefficient, especially for random access to non-contiguous blocks in the file.

**Disk Fragmentation:** With linked block allocation, small files may be scattered across the disk, leading to fragmentation. This can result in wasted space and reduced performance, as the disk needs to seek between non-contiguous blocks to read and write files.

**Limited Concurrency:** With linked block allocation, concurrent access to the same file can be difficult, especially if multiple processes are trying to modify the same block in the file. This can result in data inconsistencies and reduced performance.

**Security Vulnerabilities:** Linked block allocation can also be vulnerable to security attacks, such as buffer overflow attacks, in which an attacker can modify the pointers in a file to point to malicious code. This can allow the attacker to execute arbitrary code on the system or gain unauthorized access to sensitive data.

**Disk Space Overhead:** Linked block allocation also requires additional disk space to store the pointers between blocks, which can reduce the overall disk space utilization and waste disk space.

List any 3 problems of the indexed block allocation schemes

The indexed block allocation scheme is a file allocation scheme in which files are stored on a disk using an index block, which contains pointers to the blocks that make up the file. While this scheme has some advantages, it also has several problems, including:

**Limited File Size:** With indexed block allocation, the maximum file size is limited by the size of the index block. If the index block is full, it may not be possible to allocate additional blocks to the file, which can limit the maximum file size.

**Disk Fragmentation:** With indexed block allocation, small files may be scattered across the disk, leading to fragmentation. This can result in wasted space and reduced performance, as the disk needs to seek between non-contiguous blocks to read and write files.

**Increased Overhead:** With indexed block allocation, each file requires an additional index block to store the pointers to the file's data blocks. This can increase the amount of disk space required to store a file, as well as the amount of disk I/O operations required to read and write the file.

**Poor Performance for Random Access:** With indexed block allocation, accessing a specific block in the file requires following a chain of pointers from the index block to the desired block. This can be slow and inefficient, especially for random access to non-contiguous blocks in the file.

**Limited Concurrency:** With indexed block allocation, concurrent access to the same file can be difficult, especially if multiple processes are trying to modify the same block in the file. This can result in data inconsistencies and reduced performance.

**Disk Space Overhead:** Indexed block allocation also requires additional disk space to store the index block, which can reduce the overall disk space utilization and waste disk space.

What is a device driver? Write some 7-8 points that correctly describe need, use, placement, particularities of device drivers.

A device driver is a software program that enables communication between a computer's operating system and a hardware device. The driver acts as a translator, converting commands from the operating system into a language that the device can understand, and translating data from the device into a format that the operating system can use.

The need for a device driver arises because hardware devices have their own unique language and protocols that are specific to their function and manufacturer. The operating system is designed to communicate with devices in a generic way, so it needs a driver to translate the commands and data into the device's specific language.

The use of device drivers allows the operating system to communicate with a wide variety of hardware devices, including printers, scanners, cameras, network adapters, and many others. Without device drivers, the operating system would not be able to recognize or communicate with these devices, making them useless.

Device drivers are typically placed between the operating system and the hardware device they are designed to control. When the operating system needs to communicate with a device, it sends commands to the driver, which then sends the commands to the device in the correct format. When the device responds with data or status information, the driver translates the information back into a format that the operating system can understand.

Device drivers can be built into the operating system, or they can be separate programs that are loaded when the hardware device is connected or installed. Drivers can be updated or replaced to improve performance or fix bugs, and they are an important part of maintaining a stable and functional computer system.

What is a zombie process? How is it different from an orphan process?

-> Zombie processes are the one which are terminated but are waiting for parent to call wait() so that the process will be released from the process table, in case of orphan process parent does not invoke wait() and terminates.

A zombie process is a type of process in a computer system that has completed its execution but still has an entry in the process table. This occurs when the process has sent a message to its parent process to indicate that it has finished executing, but the parent process has not yet read the message to reap its child process. As a result, the process is in a "zombie" state, and its process ID (PID) and other resources are still allocated in the system, but the process itself is not executing.

An orphan process, on the other hand, is a type of process that is still running but has lost its parent process. This can occur when the parent process terminates or crashes before the child process completes its execution. In this case, the child process is re-parented to the init process, which becomes its new parent process. The orphan process continues to run, but it may not be able to access certain resources or perform certain operations without its original parent process.

The main difference between a zombie process and an orphan process is that a zombie process has completed its execution and is waiting for its parent process to reap it, while an orphan process is still running but has lost its parent process. Both types of processes can cause issues in a computer system, as they can consume system resources and cause performance issues. However, zombie processes are generally less problematic than orphan processes, as they are not actively executing and do not consume CPU time or other resources.

Which state changes are possible for a process, which changes are not possible?

->NEW READY RUNNING WAITING TERMINATING. That diagram of states that much possible

A process in a computer system can go through several states during its lifetime. The possible states that a process can be in include:

New: The process is being created, and its resources are being allocated.

Ready: The process is waiting to be assigned to a processor for execution.

Running: The process is being executed by a processor.

Waiting: The process is waiting for a certain event or resource, such as input/output or a lock.

Terminated: The process has finished executing and has been removed from the system.

Some state changes that are possible for a process include:

New to Ready: The process has been created and is now waiting to be assigned to a processor for execution.

Ready to Running: The process has been assigned to a processor and is now being executed.

Running to Waiting: The process is waiting for a certain event or resource.

Running to Terminated: The process has finished executing and has been removed from the system.

However, some state changes are not possible or allowed for a process, such as:

Running to Ready: A running process cannot go back to the ready state without first waiting for a certain event or resource.

Waiting to Running: A waiting process cannot resume execution without first receiving the event or resource it was waiting for.

Terminated to any other state: Once a process has finished executing and has been removed from the system, it cannot go back to any other state.

It's important to note that the state changes that are possible for a process may vary depending on the operating system and its implementation.

Top of Form

What is mkfs? what does it do? what are different options to mkfs and what do they mean?

When a storage device is formatted with a file system, the operating system can read and write data to the device, organize the data into files and directories, and track where each file is stored on the device. mkfs creates the necessary structures and metadata for the file system, such as the superblock, inode table, and data blocks.

Some of the different options available with the mkfs command and their meanings are:

-t: Specifies the type of file system to create, such as ext2, ext3, ext4, NTFS, FAT32, etc.

-c: Checks the device for bad blocks and marks them as unusable.

-b: Specifies the block size of the file system. The default value depends on the file system type.

-i: Specifies the number of inodes per block group or the inode size.

-L: Labels the file system with a name or a volume label.

-F: Forces the creation of the file system without prompting for confirmation.

-n: Displays the actions that would be taken but does not actually create the file system.

Make file system. mkfs is used to build a Linux file system on a device, usually a hard disk partition. On other operating systems, creating a file system is called *formatting*.

What is the purpose of the PCB? which are the necessary fields in the PCB?

->Process control block. Has pc,cpu register,cpu scheduling info,memory management info(base+limit), accounting info,i/o info

The PCB (Process Control Block) is a data structure used by operating systems to store information about a running process. The PCB is an essential component of a process management system, as it provides the operating system with the necessary information to manage and control processes effectively.

The PCB contains a wide range of information about a process, including:

Process State: The current state of the process, such as running, ready, waiting, etc.

Process ID: A unique identifier assigned to the process by the operating system.

Program Counter: The address of the next instruction to be executed.

CPU Registers: The values of CPU registers when the process was last interrupted.

Memory Management Information: Information about the memory allocated to the process, including its base address, size, and access permissions.

I/O Status Information: The status of any open input or output requests made by the process.

Accounting Information: Information about the resources used by the process, including CPU time, memory usage, and I/O operations.

The necessary fields in a PCB may vary depending on the operating system and its implementation. However, some of the essential fields that are typically included in a PCB are:

Process ID: A unique identifier assigned to the process by the operating system.

Process State: The current state of the process, such as running, ready, waiting, etc.

Program Counter: The address of the next instruction to be executed.

CPU Registers: The values of CPU registers when the process was last interrupted.

Memory Management Information: Information about the memory allocated to the process, including its base address, size, and access permissions.

I/O Status Information: The status of any open input or output requests made by the process.

The PCB is a critical data structure that allows the operating system to efficiently manage and control processes in a multi-tasking environment.

Write a C program that uses globals, local variables, static local variables, static global variables, and malloced memory. Compile it. Dump the object code file using `objdump -d -x`. Can you see in the output, the separation into stack, heap, text, etc?

Which parts of a C program ((typedef, #define, #include, functions, local vars, statics, globals, #ifdef, ... ) etc occupy RAM when it's a process, and where (code, data, stack, ..)?

code- main code function code

data-statics, globals,

stack - local vars

heap- mallocated memory

#define, #include – no memory

typedef: A typedef declaration does not occupy any memory at runtime, as it is only used by the compiler to define new types.

#define: Similarly, a #define statement does not occupy any memory at runtime, as it is only used by the preprocessor to replace text in the code before compilation.

#include: An #include statement brings in code from another file, which may contain any of the other components listed here. When the program is compiled, the contents of the included file are merged into the main program's code section.

Functions: The code for each function in a C program is stored in the code section of memory, which is typically read-only.

Local variables: Local variables are created on the stack when a function is called, and are destroyed when the function returns. They are stored in the stack section of memory.

Static variables: Static variables are initialized at program startup and persist throughout the lifetime of the program. They are typically stored in the data section of memory.

Global variables: Global variables are also initialized at program startup and persist throughout the lifetime of the program. They are typically stored in the data section of memory.

#ifdef: #ifdef statements are used by the preprocessor to conditionally include or exclude code based on whether a certain macro is defined. The resulting code is merged into the main program's code section.

Describe the role of CPU (MMU), kernel, compiler in the overall task of "memory management" for a process/program.

Memory management is a critical aspect of running a process or program, and it involves the cooperation of several components, including the CPU's memory management unit (MMU), the operating system kernel, and the compiler used to build the program.

The role of the CPU's memory management unit (MMU) is to translate virtual memory addresses used by the process into physical memory addresses that correspond to locations in RAM. The MMU does this by using a hardware mechanism called memory mapping, which maps the virtual memory space of the process to a physical address space in RAM. The MMU also provides hardware-based memory protection by preventing a process from accessing memory outside of its allocated virtual address space, which helps to prevent errors and security vulnerabilities.

The kernel of the operating system is responsible for managing the allocation and deallocation of physical memory for all processes running on the system. The kernel maintains a mapping of physical memory pages to virtual memory pages used by each process, and it performs operations such as paging, swapping, and allocating new memory pages as needed. The kernel also provides mechanisms for inter-process communication and synchronization, which may involve sharing memory between processes.

The compiler used to build a program plays a critical role in memory management as well. The compiler determines the layout of the program's code, data, and stack in memory, and it generates machine code that interacts with the MMU and the kernel to allocate and access memory. The compiler may use techniques such as static memory allocation, dynamic memory allocation, and garbage collection to manage the program's memory usage, depending on the programming language and the program's requirements.

In summary, memory management for a process or program is a complex task that requires the cooperation of the CPU's MMU, the operating system kernel, and the compiler used to build the program. The MMU translates virtual memory addresses to physical memory addresses and provides memory protection, the kernel manages the allocation and deallocation of physical memory for all processes, and the compiler determines the layout of the program's memory usage and generates machine code that interacts with the MMU and kernel to allocate and access memory.

What is the difference between a named pipe and un-named pipe?

-> unnamed pipe/ ordinary pipe is only used for communication between a child and parent process, while a named pipe there is no need of child and parent communication, follows fifo. Processes of different ancestry can share data through a named pipe. bidirectional. unnamed = unidirectional

named = pipe act as file

"unnamed" and lasts only as long as the process. A named pipe, however, can last as long as the system is up, beyond the life of the process. It can be deleted if no longer used

A named pipe and an unnamed pipe are two types of interprocess communication (IPC) mechanisms available on Unix-like operating systems. The main difference between these two types of pipes is in how they are created and accessed.

An unnamed pipe, also known as a regular pipe, is a type of IPC that allows data to be transferred between two related processes, typically a parent and a child process. The pipe is created by calling the pipe system call, which creates a pipe buffer in memory that is used to transfer data between the processes. An unnamed pipe has no name or file descriptor associated with it and is accessible only to the related processes that share the pipe.

A named pipe, also known as a FIFO (first-in, first-out) or a named FIFO, is a type of IPC that allows data to be transferred between any two or more unrelated processes on the system. A named pipe is created by calling the mkfifo system call, which creates a named file in the file system that acts as a pipe. The named pipe has a name and is represented by a file descriptor that can be used by any process with appropriate permission to access the pipe.

In summary, the main difference between a named pipe and an unnamed pipe is that a named pipe is accessible to any process on the system that has permission to access it, while an unnamed pipe is accessible only to related processes that share the pipe. A named pipe is created using a file system file, while an unnamed pipe is created using a pipe buffer in memory.

Describe the steps involved in resolving the path name /a/b/c on an ext2 filesystem.

The path name `/a/b/c` is broken down into its individual components, separated by the directory separator `/`.

Starting from the root directory `/`, the ext2 filesystem is searched for the first component `a`. The root directory is always represented by inode number 2 on an ext2 filesystem.

The directory entry for `a` is located in the root directory, which contains an inode number pointing to the inode of directory `a`.

The inode of directory `a` is read from disk, and its contents are parsed to find the next component `b`.

The directory entry for `b` is located in directory `a`, which contains an inode number pointing to the inode of directory `b`.

The inode of directory `b` is read from disk, and its contents are parsed to find the final component `c`.

The directory entry for `c` is located in directory `b`, which contains an inode number pointing to the inode of file or directory `c`.

The inode of file or directory `c` is read from disk, and its type and permissions are checked to ensure that the process has the necessary access to the file or directory.

If `c` is a directory, the process may repeat steps 4-8 for any further components in the path name.

Once the final component has been resolved, the process can access the file or directory `c` using its inode number and perform any desired operations.

In summary, resolving a path name on an ext2 filesystem involves recursively traversing the directory tree to locate each component of the path name, until the final file or directory is found and its inode can be used to access its contents.

Suppose there is an ext2 partition with block size of 2KB. what is the maximum possible size of a file in kilobytes?

The maximum possible size of a file on an ext2 partition depends on the block size used by the filesystem.

In this case, the block size is 2KB, which means that each file can occupy a maximum number of blocks equal to the total size of the file divided by the block size.

The maximum number of blocks that can be addressed by an ext2 filesystem with a block size of 2KB is  $2^{32} - 1$ , or about 4 billion blocks. This means that the maximum possible size of a file in kilobytes can be calculated as follows:

Maximum file size = (Maximum number of blocks \* Block size) / 1024 = (4 billion \* 2KB) / 1024 = 8,388,608 KB

Therefore, the maximum possible size of a file on an ext2 partition with a block size of 2KB is 8,388,608 KB, or approximately 8.4 GB.

Write a program that implements a pipe between two processes only (this is Shell assignment-2, first small part).

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
```

```
int main(void) {
    int fd[2];
    pid_t pid;

    char buf[1024];

    // Create the pipe
    if (pipe(fd) == -1) {
        perror("pipe");
        exit(EXIT_FAILURE);
    }
```

```

// Fork a child process
pid = fork();

if (pid == -1) {
    perror("fork");
    exit(EXIT_FAILURE);
}

if (pid == 0) { // Child process
    close(fd[1]); // Close the write end of the pipe

    while (read(fd[0], buf, sizeof(buf)) > 0) { // Read from the read end of the pipe
        printf("Child received message: %s", buf);
    }

    close(fd[0]); // Close the read end of the pipe
    exit(EXIT_SUCCESS);
} else { // Parent process
    close(fd[0]); // Close the read end of the pipe

    char *msg = "Hello, child process!\n";
    write(fd[1], msg, strlen(msg)); // Write to the write end of the pipe

    close(fd[1]); // Close the write end of the pipe
    exit(EXIT_SUCCESS);
}
}

```

What are privileged instructions? Why are they required? What is their relationship to 2 modes of CPU execution?

2 modes kernel mode and user mode . Privileged instruction only run by kernel.kernel=0

user mode bit=1.protection coz of dualmode

Privileged instructions are CPU instructions that can only be executed in privileged mode. These instructions are designed to perform operations that are critical to the functioning of the operating system, such as manipulating hardware registers, controlling system resources, and accessing privileged data structures. Examples of privileged instructions include loading or modifying control registers, executing I/O operations, and changing the memory protection level.

Privileged instructions are required to ensure the safety and stability of the operating system. By restricting access to certain CPU instructions and system resources, the operating system can prevent user-level processes from interfering with critical system functions and compromising the security of the system. For example, if a user-level process were able to execute privileged instructions directly, it could potentially cause the system to crash or allow malicious actors to gain control of the system.

Most CPUs support two modes of execution: user mode and privileged mode (also known as kernel mode). User mode is the mode in which most applications and processes run, and it provides restricted access to system resources and privileged instructions. Privileged mode, on the other hand, is the mode in which the operating system kernel runs, and it provides full access to system resources and privileged instructions. In privileged mode, the CPU can execute privileged instructions and access system resources that are restricted in user mode. By contrast, in user mode, the CPU is limited to executing a subset of instructions that do not require privileged access. This separation of privileges between user mode and privileged mode helps to ensure the security and stability of the system by preventing user-level processes from directly accessing sensitive system functions and resources.



Explain what happens in pre-processing, compilation and linking and loading.

In general, the process of compiling and linking a program involves several stages: pre-processing, compilation, linking, and loading. Here's a brief explanation of what happens at each stage:

**Pre-processing:** During this stage, the pre-processor takes the source code of a program and performs various operations on it, such as including header files, replacing macros with their definitions, and removing comments. The output of this stage is a modified version of the original source code, which is then passed on to the compiler.

**Compilation:** During the compilation stage, the compiler takes the pre-processed source code and translates it into machine code, which is the binary code that can be executed by the CPU. This involves several sub-stages, including lexical analysis (breaking the code into tokens), syntax analysis (parsing the tokens into a syntax tree), semantic analysis (checking the syntax tree for correctness), code generation (generating machine code from the syntax tree), and optimization (improving the efficiency of the generated code). The output of this stage is an object file, which contains the machine code for a single compilation unit.

**Linking:** During the linking stage, the linker takes one or more object files and combines them into a single executable file or shared library. This involves several sub-stages, including symbol resolution (matching up function and variable references across different object files), relocation (adjusting the machine code of the object files to work together), and library linking (linking in external libraries that are required by the program). The output of this stage is an executable file or shared library, which can be executed or loaded by the operating system.

**Loading:** During the loading stage, the operating system loads the executable file or shared library into memory and prepares it for execution. This involves several sub-stages, including memory allocation (allocating memory for the program and its data), symbol binding (mapping symbols to their memory addresses), and relocation (adjusting the machine code to work with the actual memory layout). Once the loading is complete, the program can be executed by the CPU.

Overall, the process of compiling and linking a program involves several complex stages, each of which is responsible for a different aspect of the process. By breaking the process down into these stages, compilers and linkers can perform each stage more efficiently and accurately, leading to faster and more reliable program builds.

### **What does lgdt do?**

The LGDT instruction is used in x86 architecture to load the Global Descriptor Table (GDT) register. The GDT is a data structure used by the processor to define the various memory segments used by the operating system and programs running on the computer.

When the LGDT instruction is executed, it loads the address of the GDT into the GDTR (Global Descriptor Table Register), which tells the processor where the GDT is located in memory. The GDT contains a set of descriptors that define the various memory segments used by the system, including code segments, data segments, stack segments, and task state segments.

By loading the GDT register using the LGDT instruction, the processor can then use the descriptors in the GDT to properly manage memory and execute instructions in different memory segments as needed. Overall, the LGDT instruction is an important part of managing system memory and ensuring proper operation of the operating system and applications running on a computer.

### **Where is the initial GDT specified and where?**

The initial GDT (Global Descriptor Table) is typically specified and defined by the operating system during system initialization. The location of the GDT is determined by the operating system, and it is usually stored in a reserved area of memory known as the system data area or system descriptor table.

During the boot process, the BIOS initializes the processor and sets up the initial memory environment, including the initial GDT. Once the operating system takes control, it can modify the GDT as needed to define memory segments and manage system memory.

The GDT is a data structure used by the x86 architecture to define the various memory segments used by the processor. It consists of a set of descriptors, which define the base address, limit, and access permissions of each memory segment. These descriptors are typically created by the operating system and stored in the GDT for use by the processor.

In summary, the initial GDT is specified and defined by the operating system during system initialization and is usually stored in a reserved area of memory known as the system data area or system descriptor table.

### **How does xv6 utilise the segmentation of x86 ?**

xv6 is a simple UNIX-like operating system designed for teaching purposes and implemented for the x86 architecture. In xv6, segmentation is used to implement memory protection, which is an essential feature of operating systems.

The x86 architecture supports two types of memory segmentation: the Global Descriptor Table (GDT) and the Local Descriptor Table (LDT). In xv6, the GDT is used to define the memory segments used by the kernel and user-level programs, and to enforce memory protection.

Specifically, xv6 defines four segments in the GDT:

- The null segment: This is a dummy segment used to represent a null or invalid segment descriptor.
- The kernel code segment: This is a segment that contains the kernel code and is accessible only in kernel mode.
- The kernel data segment: This is a segment that contains the kernel data and is accessible only in kernel mode.
- The user segment: This is a segment that contains the user-level program code and data and is accessible only in user mode.

Each of these segments is defined using a descriptor in the GDT, which specifies the base address, size, access permissions, and other attributes of the segment. When a program executes, the processor uses these descriptors to determine whether a memory access is valid or not, and to enforce memory protection.

In summary, xv6 uses segmentation in the x86 architecture to implement memory protection by defining memory segments using the Global Descriptor Table (GDT). This allows xv6 to enforce access permissions and prevent programs from accessing memory they are not authorized to access.

### **What is the meaning of "selector" in x86 and how is it fixated in xv6 code?**

In x86 architecture, a selector is a 16-bit value that is used to select a descriptor from the Global Descriptor Table (GDT) or Local Descriptor Table (LDT). A selector consists of three fields: the segment index, the request privilege level (RPL), and the table indicator.

The segment index is an index into the GDT or LDT that identifies the descriptor associated with the selector. The RPL specifies the privilege level of the code that is attempting to access the segment, and the table indicator specifies whether the selector refers to the GDT or LDT.

In xv6, selectors are used extensively to access memory segments and enforce memory protection. For example, when a user-level process attempts to access a memory location, the processor checks the selector associated with the memory access to ensure that the process has the proper access permissions.

In xv6, the selectors are fixated or set during the initialization of the kernel. Specifically, during system boot, xv6 sets up the GDT and defines the selectors for the kernel code, kernel

data, and user code and data segments. These selectors are then used throughout the operating system to access the various memory segments.

Overall, selectors are a critical part of memory segmentation in x86 architecture and are used by xv6 to implement memory protection and manage memory access for the kernel and user-level processes.

### **How do you enable 4MB pages?**

Enabling 4MB pages on x86 architecture involves setting a bit in the page directory entry for the desired page directory entry. Here are the steps to enable 4MB pages:

Set the PSE (Page Size Extension) bit in the CR4 (Control Register 4) register. This bit enables support for 4MB pages.

Allocate a page directory entry for the desired page directory entry.

Set the P (present) and PS (page size) bits in the page directory entry. The P bit indicates that the page is present in memory, and the PS bit indicates that the page size is 4MB.

Set the physical base address of the page in the page directory entry.

Repeat steps 2-4 for each page that should be mapped to 4MB pages.

Once the page directory entries are set up, the operating system can use 4MB pages to improve performance by reducing the number of page table entries required to map a large amount of memory.

It's important to note that enabling 4MB pages may not always be beneficial, as it can increase memory fragmentation and decrease flexibility in memory management. It's best to evaluate the specific use case and system requirements before enabling 4MB pages.

Overall, enabling 4MB pages involves setting a bit in the CR4 register and setting the appropriate bits in the page directory entries for the desired pages.

### **What are the 3 segments in 'kernel' file?**

The three segments in the kernel file of an operating system, specifically in the context of xv6, are:

- The text segment: This segment contains the executable code of the kernel, including functions, instructions, and data tables that are needed for the kernel to operate. This segment is usually read-only to prevent accidental modification of the kernel code.
- The data segment: This segment contains initialized data used by the kernel, such as global variables and static variables. This segment is usually read-write to allow the kernel to modify the data as needed.
- The bss segment: This segment contains uninitialized data used by the kernel, such as uninitialized global variables and static variables. This segment is also read-write and is initialized to zero during system startup.

These three segments are defined in the kernel.ld linker script file in xv6. The linker script specifies the starting addresses and sizes of the three segments in the kernel file, as well as other linker options such as memory layout and symbol resolution.

In summary, the three segments in the kernel file of xv6 are the text segment, the data segment, and the bss segment. These segments contain executable code, initialized data, and uninitialized data used by the kernel to operate.

### **At which Virtual address and which physical address the code of kernel loaded?**

The virtual and physical addresses where the kernel code is loaded can vary depending on the specific operating system and hardware platform. However, in the case of xv6 running on the QEMU emulator, the kernel is loaded at virtual address 0x80000000 and physical address 0x00100000.

During the boot process, the bootloader (bootasm.S) loads the kernel image from the disk into memory, starting at physical address 0x00100000. The bootloader sets up the page tables and maps the kernel image into the virtual address space at address 0x80000000.

This mapping is defined in the kernel.ld linker script, which specifies the starting virtual address and the physical address where the kernel is loaded. The linker script sets the kernel base virtual address to 0x80000000 and the physical address to 0x00100000.

Once the kernel is loaded into memory and mapped to the virtual address space, the kernel code can be executed and the operating system can begin functioning.

### **Can you suggest an alternative GDT to make the boot process still work?**

The choice of a Global Descriptor Table (GDT) is dependent on the specific requirements of the operating system and hardware platform. However, there are alternative GDT configurations that can be used to make the boot process still work.

One alternative is to use a flat GDT, which consists of a single descriptor that covers the entire address space. This configuration simplifies the GDT setup and management, but it may not provide the same level of memory protection and segmentation as a more complex GDT.

Another alternative is to use a segmented GDT with a smaller number of descriptors than the default GDT used by most operating systems. This can help reduce the size of the GDT and simplify the setup process, while still providing some level of memory protection and segmentation.

It's important to note that choosing an alternative GDT configuration may impact the performance and security of the operating system. Therefore, it's important to carefully evaluate the specific requirements and constraints of the system before making any changes to the GDT configuration. Additionally, any modifications to the GDT should be thoroughly tested to ensure that they do not cause any issues during the boot process or normal operation of the operating system.

### **What is the size of ELF header?**

The size of the ELF header depends on the class (32-bit or 64-bit) and the encoding (little-endian or big-endian) of the ELF file.

For a 32-bit little-endian ELF file, the size of the ELF header is 52 bytes.

For a 64-bit little-endian ELF file, the size of the ELF header is 64 bytes.

For a big-endian ELF file, the size of the ELF header is the same as for the little-endian ELF file of the same class.

It's important to note that the ELF header is only the first part of the ELF file format. It contains information about the file's architecture, type, entry point, and other attributes, but it does not contain the actual program code or data. The rest of the ELF file consists of program headers (describing segments of the program) and section headers (describing sections of the program). Together, these components make up the complete ELF file format.

### **Why is 4K data being read initially in bootmain()?**

In the xv6 operating system, the bootmain() function is responsible for loading the kernel image from the disk into memory and preparing the system for the kernel to start executing.

The reason that bootmain() reads 4K (4096 bytes) of data initially is because it needs to read the first block of the kernel image from the disk. This block contains the initial boot loader code (in bootasm.S) that loads the rest of the kernel image.

In the `bootmain()` function, the initial block of the kernel image is read from the disk into memory using the `readsect()` function. This function reads one sector (512 bytes) at a time from the disk, so `bootmain()` reads 8 sectors (4096 bytes) to read the first 4K block of the kernel image.

Once the initial block of the kernel image is loaded, `bootmain()` passes control to the `kernel_entry()` function, which starts executing the kernel code. The rest of the kernel image is loaded and set up by the `kernel_main()` function.

Therefore, reading 4K of data initially is a necessary step in the boot process of xv6 to load the initial boot loader code and start the kernel.

**The kernel is expected to be in sector-1. Isn't it that `bootmain()` is actually reading it from sector-0? How is that ?**

You are correct that in the xv6 operating system, the kernel is expected to be located in sector 1 of the disk, not sector 0. However, in the `bootmain()` function, the kernel image is actually read from sector 0 of the disk, not sector 1.

This is because the initial bootloader code (in `bootasm.S`) that is located in sector 0 of the disk is responsible for loading the rest of the kernel image into memory. Therefore, the first sector that `bootmain()` reads is actually sector 0, which contains the bootloader code that loads the rest of the kernel from sector 1 onwards.

The bootloader code in `bootasm.S` is designed to be small enough to fit within a single disk sector. This code sets up the processor and memory, reads the kernel image from the disk into memory, and passes control to the kernel entry point. After the bootloader code has finished its job, it is no longer needed and is overwritten by the kernel image as it is loaded into memory.

Therefore, although the kernel is expected to be in sector 1 of the disk, `bootmain()` actually reads it from sector 0 via the initial bootloader code.

**Why do you think the paging is enabled later, after `bootmain()` is over?**

In the xv6 operating system, the `bootmain()` function is responsible for loading the kernel image from disk and setting up the initial state of the system. However, the system is not fully initialized at this point, and there are many tasks that still need to be performed before the kernel is ready to start executing user programs.

One of these tasks is setting up the memory management unit (MMU) to enable paging. Paging is a mechanism used by the operating system to manage physical memory, allowing the system to map virtual addresses used by user programs to physical addresses in memory.

The reason why paging is enabled later, after `bootmain()` is over, is because setting up paging requires more resources and is more complex than simply loading the kernel image into memory. In particular, enabling paging requires setting up page tables, which are data structures used by the MMU to map virtual addresses to physical addresses. Setting up the page tables requires knowledge of the layout of physical memory, as well as the virtual address space used by the kernel and user programs.

Therefore, it is more practical to set up paging after `bootmain()` is over, when the system is in a more fully initialized state and there is more information available about the layout of physical memory and the virtual address space used by the system. This allows for more accurate and efficient mapping of virtual addresses to physical addresses, which can improve system performance and stability.

**Where is the "end" variable, the first argument to `knit1()` declared? What's its value?**

In the xv6 operating system, the `end` variable is declared in the kernel linker script (`kernel.ld`) and is defined as the address of the end of the kernel image in memory.

The kernel.ld linker script is used to specify the layout of the kernel image in memory, including the location of different sections of the kernel code and data. In this script, the end symbol is defined using the following line:

```
scss
```

```
Copy code
```

```
/* End of kernel image */
```

```
. = ALIGN(4096);
```

```
end = .;
```

This line sets the end symbol to the current address in memory (denoted by the . symbol) after it has been aligned to a 4096-byte boundary using the ALIGN directive. This ensures that the end symbol is aligned to a page boundary, which is necessary for memory management

The end symbol is passed as the first argument to the kinit1() function, which is called by main() to initialize the kernel memory allocator. The kinit1() function uses the end symbol to determine the starting address of the free memory region, which it then uses to allocate memory for the kernel data structures.

The value of the end symbol is determined at compile-time based on the size of the kernel image and the layout specified in kernel.ld. The exact value of end depends on the specific version of xv6 being used, as well as the configuration of the system on which it is running. However, it is typically set to a value that is just below the top of physical memory, leaving the rest of the memory for use by user programs.

### **The PGROUNDUP does an adjustment on a page boundary, that is ceiling or floor?**

The PGROUNDUP macro in xv6 adjusts a given address to the next higher multiple of the page size, which is 4096 bytes in xv6. This means that it rounds the given address up to the next page boundary.

For example, if addr is an address that is not aligned to a page boundary, the expression PGROUNDUP(addr) will evaluate to the address of the next page boundary after addr.

In contrast, the PGROUNDDOWN macro in xv6 adjusts a given address to the previous lower multiple of the page size. This means that it rounds the given address down to the previous page boundary.

For example, if addr is an address that is not aligned to a page boundary, the expression PGROUNDDOWN(addr) will evaluate to the address of the previous page boundary before addr.

Therefore, PGROUNDUP rounds up to the ceiling of the page boundary, while PGROUNDDOWN rounds down to the floor of the page boundary.

### **The free list maintained by xv6 is which type of a linked-list?**

In xv6, the free list that is used to manage physical memory pages is implemented as a singly-linked list. Each node in the list represents a free page of memory, and contains a pointer to the next free page in the list.

The free list is maintained by the kernel's memory allocator, which keeps track of the free pages of memory and allocates them as needed for kernel data structures, user processes, and other purposes. When a page of memory is no longer needed, it is returned to the free list, which makes it available for future allocations.



The struct run data structure is used to represent a node in the free list. It contains a single field, struct run \*next, which is a pointer to the next node in the list. When a page is added to the free list, it is inserted at the beginning of the list, and its next field is set to the previous head of the list.

Because the free list is implemented as a singly-linked list, iterating over the list requires traversing each node in sequence, starting from the head of the list. This can be less efficient than a doubly-linked list, which allows for faster iteration in both directions. However, the use of a singly-linked list simplifies the implementation of the free list and reduces the amount of memory overhead required to maintain the list.

### **What's the maximum possible value that PHYSTOP can be set to ?**

In xv6, the PHYSTOP constant is defined in the file param.h as the maximum physical memory address that the kernel can use. The maximum value of PHYSTOP is determined by the size of physical memory on the system.

By default, PHYSTOP is set to 0x800000 (8 megabytes) in the xv6 codebase. However, this value can be changed depending on the amount of physical memory available on the system.

In general, the maximum value of PHYSTOP will be determined by the amount of physical memory that the system can address using its hardware architecture. For a 32-bit system, the maximum amount of physical memory that can be addressed is 4 gigabytes, which corresponds to a value of 0xFFFFFFFF in hexadecimal. However, not all of this memory may be available for use by the kernel, as some memory may be reserved for hardware devices or other purposes. For a 64-bit system, the maximum amount of physical memory that can be addressed is much higher, up to several terabytes or more, depending on the hardware architecture.

### **kalloc returns which free frame? (first/last/best/LIFO/FIFO/LFU/LRU/etc? )**

In xv6, the kalloc function is responsible for allocating a single physical memory page (4KB) from the free list maintained by the kernel's memory allocator. The kalloc function selects the next available page from the free list using a first-fit strategy.

In a first-fit strategy, the memory allocator scans the free list from the beginning, and selects the first available memory block that is large enough to satisfy the allocation request. This strategy is relatively simple and efficient, but may result in memory fragmentation if small gaps are left between allocated blocks.

Once kalloc finds a free memory block that is large enough to satisfy the allocation request, it removes the block from the free list and returns a pointer to the start of the block as the allocated page. The function also clears the contents of the allocated page to ensure that it contains zeros.

Overall, the free list maintained by the kernel's memory allocator is implemented as a singly-linked list of free memory blocks, with each block represented by a struct run node. The kalloc function uses a first-fit strategy to select the next available memory block from the free list, and removes the block from the list once it has been allocated.

### **Describe in least possible words, the work done by kvmalloc.**

In xv6, `kvmalloc` is a function that allocates a contiguous block of kernel virtual memory (KVA) of a specified size, using the kernel's memory allocator. It returns a pointer to the start of the allocated memory block, or `NULL` if the allocation fails. Unlike `kalloc`, `kvmalloc` does not allocate physical memory pages directly, but instead allocates virtual memory pages that are backed by physical memory pages as needed. The virtual memory pages are mapped into the kernel's address space using the paging mechanism of the x86 processor, and are managed using a combination of the kernel's memory allocator and the x86 hardware memory management unit (MMU).

### **"kmap" stands for what?**

In xv6, `kmap` is a function that maps a physical memory page to a kernel virtual address (KVA) in the kernel's address space. The `kmap` function is used to create a temporary mapping between a physical memory page and a KVA, allowing the kernel to access the contents of the page using normal memory operations.

The name `kmap` is short for "kernel map", which refers to the process of mapping physical memory to virtual memory in the kernel's address space. By mapping physical memory to a KVA, the kernel can manipulate the contents of the memory page using normal memory operations, such as reading and writing.

The `kmap` function is typically used in conjunction with `kunmap`, which is used to unmap a previously-mapped physical memory page from the kernel's address space. The `kmap` and `kunmap` functions are used to manage temporary mappings between physical and virtual memory in the kernel's address space, allowing the kernel to access the contents of physical memory pages as needed.

### **Which function does mappings in Page table ? (not page directory)**

In xv6, the function that creates and modifies page table entries (PTEs) is called `mappages()`. The `mappages()` function takes as arguments a page directory pointer (`pgdir`), a virtual address (`va`), a size (`size`), and a physical address (`pa`), and creates or modifies PTEs in the page table(s) that map the virtual address range [`va`, `va+size`) to the physical address range [`pa`, `pa+size`).

The `mappages()` function uses the x86 paging mechanism to create or modify PTEs in the page table(s) associated with the specified page directory. The function sets the appropriate page table flags and permissions to ensure that the mapping is valid and secure, and updates the page directory entry (PDE) to point to the new page table(s) if necessary.

Overall, the `mappages()` function is an important part of the xv6 kernel's memory management system, and is used to create and manage virtual-to-physical memory mappings in the kernel's address space.

### **Which function does mappings in Page directory?**

In xv6, the function that creates and modifies page directory entries (PDEs) is called `setupkvm()`. The `setupkvm()` function is responsible for setting up the kernel's page directory and initial kernel virtual-to-physical memory mappings during the kernel boot process.



The `setupkvm()` function creates the kernel's page directory and maps the first PHYSTOP physical memory pages to the corresponding kernel virtual addresses, using a combination of `mappages()` and `kmap()`. The function sets the appropriate PDE flags and permissions to ensure that the mappings are valid and secure, and initializes the kernel's free memory list using the remaining physical memory pages.

Overall, the `setupkvm()` function is an important part of the xv6 kernel's memory management system, and is used to set up the initial kernel virtual-to-physical memory mappings and free memory list during the boot process.

### **Which of the flags in PTE and PDE are actually used in xv6 and which flags are ignored?**

In xv6, the page table entry (PTE) and page directory entry (PDE) flags that are used and ignored depend on the specific context in which they are used. Here is a brief overview:

Used flags in PTE:

- PTE\_P (present bit): used to indicate whether the page is present in memory or not.
- PTE\_W (writable bit): used to indicate whether the page is writable or not.
- PTE\_U (user bit): used to indicate whether the page is accessible in user mode or not.
- PTE\_PG (page table entry bit): used to indicate whether the PTE represents a page table or a physical page.
- Ignored flags in PTE:
- PTE\_AVAIL: these bits are reserved for software use and are not used by xv6.

Used flags in PDE:

- PTE\_P (present bit): used to indicate whether the page table or physical page is present in memory or not.
- PTE\_W (writable bit): used to indicate whether the page table or physical page is writable or not.
- PTE\_U (user bit): used to indicate whether the page table or physical page is accessible in user mode or not.

Ignored flags in PDE:

- PTE\_PS (page size bit): xv6 does not use large 2MB or 4MB pages, so this bit is ignored.
- PTE\_A (accessed bit): xv6 does not use the accessed bit for page tables or physical pages, so this bit is ignored.
- PTE\_D (dirty bit): xv6 does not use the dirty bit for page tables or physical pages, so this bit is ignored.
- PTE\_AVAIL: these bits are reserved for software use and are not used by xv6.

Overall, the PTE and PDE flags used and ignored by xv6 reflect the specific needs and design choices of the xv6 kernel's memory management system.

### **Why do you think `memset()` uses the `stosl/stosb` machine instructions, and does not run a loop to do the copy?**

The `memset()` function in xv6 uses the `stosl` (store long) and `stosb` (store byte) instructions to set a range of memory to a particular value. These instructions store a long (32 bits) or byte (8 bits) value to a memory address and increment the address pointer by the size of the stored value.

Using `stosl` and `stosb` instructions is generally faster than running a loop to copy the values one by one. This is because these instructions take advantage of the processor's ability to do bulk memory operations and avoid the overhead of looping and branching instructions.

Furthermore, using `stosl` and `stosb` instructions is a standard optimization technique and is likely to be more efficient than a hand-written loop, especially when dealing with large amounts of memory.

Overall, the use of `stosl` and `stosb` instructions in `memset()` in `xv6` is a performance optimization that takes advantage of the processor's bulk memory operations to speed up the memory setting process.

**Describe in least possible words, the work done by `walkpgdir()`.**

`walkpgdir()` is a function in `xv6` that takes a virtual address and returns the corresponding page table entry in the page directory. It does this by recursively traversing the page directory and page tables, using the page directory to locate the appropriate page table and page table entry for the given virtual address. If the page table or page table entry does not exist, `walkpgdir()` creates them. The function returns a pointer to the page table entry for the given virtual address.

**While doing the page table mappings, using `mappages()` in which function is the actual page-frame allocated? (trick question).**

The actual page-frame allocation is not done in the `mappages()` function in `xv6`. Instead, the `mappages()` function assumes that the physical memory frame to be mapped is already allocated and simply updates the corresponding page table entry in the page directory. The allocation of physical memory frames is done by other functions such as `kalloc()` and `page_alloc()`.

The `mappages()` function in `xv6` takes a page directory, a virtual address, the size of the memory to be mapped, the physical address of the memory to be mapped, and the flags for the page table entry. The function calculates the number of pages to be mapped based on the size of the memory and calls `walkpgdir()` to get a pointer to the page table entry for the given virtual address. If the page table does not exist, `mappages()` creates it using `page_alloc()`. Finally, the function sets the page table entry to the given physical address and flags for each page in the range to be mapped.

**What's the use of `kpgdir`?**

In `xv6`, `kpgdir` is a global variable that holds the page directory for the kernel address space. It is used by the kernel to manage memory mappings for the kernel code and data. By using a separate page directory for the kernel address space, the kernel can ensure that user space and kernel space do not overlap and that user programs cannot access kernel memory.

The `kpgdir` variable is initialized during the boot process by the `mem_init()` function, which sets up the initial memory mappings for the kernel. The kernel code and data are mapped to the same virtual addresses in every process's address space, so the `kpgdir` is shared by all processes.

By using a separate page directory for the kernel address space, `xv6` provides an extra layer of protection against bugs and malicious code that could try to modify or access kernel memory. The `kpgdir` variable is an important part of `xv6`'s memory management system and helps to ensure the stability and security of the operating system.

### **In lcr3(V2P(kpgdir)); why is there a call to V2P?**

In the call `lcr3(V2P(kpgdir))`, the `V2P` macro is used to convert a virtual address to a physical address. The reason for this is that the `lcr3` instruction expects a physical address as its argument, not a virtual address.

In `xv6`, the `kpgdir` variable holds the kernel's page directory, which contains virtual-to-physical mappings for the kernel's memory. However, the `lcr3` instruction requires a physical address for the new page directory base address to be loaded into the `CR3` control register.

The `V2P` macro is used to convert the virtual address of the `kpgdir` to its corresponding physical address before passing it as an argument to `lcr3`. This ensures that the correct physical address is loaded into the `CR3` register and that the kernel's page directory is properly initialized.

In summary, the `V2P` macro is used to translate a virtual address to a physical address, which is necessary for properly initializing the `CR3` control register with the physical address of the kernel's page directory.

Questions for preparation of T1 (OS 2023)

--

**List all the blocks/data-structures on the ext2 partition (e.g. superblock, group descriptor, block bitmap, directory data block, etc) with block size 4KB, that MAY be updated (means changed) when a file of size 8 KB is created.**

1. Superblock: The superblock contains information about the file system, such as the block size, inode count, and block count. When the block size changes, the superblock must be updated to reflect the new block size.
2. Block Group Descriptor Table: The block group descriptor table contains information about each block group in the file system, including the number of blocks and inodes in each block group. When the block size changes, the block group descriptor table must be updated to reflect the new block size.
3. Inode Table: The inode table contains information about each file and directory in the file system. When the block size changes, the inode table must be updated to reflect the new block size.
4. Data Blocks: The data blocks contain the actual file and directory data. When the block size changes, the data blocks must be reorganized and possibly moved to different locations on the partition to align with the new block size.
5. Block Bitmap: The block bitmap tracks which blocks are in use and which blocks are free. When the block size changes, the block bitmap must be updated to reflect the new block size.
6. Inode Bitmap: The inode bitmap tracks which inodes are in use and which inodes are free. When the block size changes, the inode bitmap must be updated to reflect the new block size.

**List any 3 problems of each of the block allocation schemes (continuous, linked, indexed)**

The continuous block allocation scheme is a file allocation scheme in which files are stored on a disk in contiguous blocks. While this scheme has some advantages, it also has several problems, including:

1. External Fragmentation: As files are added, deleted, and modified, the free space on the disk becomes fragmented into smaller and smaller pieces, making it difficult to allocate contiguous blocks for new files. This can lead to wasted space and reduced performance.
2. Limited File Size: With continuous block allocation, the maximum file size is limited by the size of the largest contiguous block of free space on the disk. If the disk becomes heavily fragmented, it may not be possible to allocate a contiguous block large enough to hold a large file.

3. **Slow File Allocation:** With continuous block allocation, finding a large enough contiguous block of free space to allocate to a new file can be time-consuming, especially on a heavily fragmented disk. This can lead to slower file allocation times and reduced performance.
4. **Difficult to Extend Files:** With continuous block allocation, extending the size of a file can be difficult, especially if there is no contiguous free space immediately following the file's current location. This may require moving the entire file to a new location on the disk, which can be time-consuming and may require significant disk I/O operations.
5. **Disk Defragmentation Required:** To overcome the issue of external fragmentation, the disk needs to be defragmented frequently. This process can be time-consuming and may require a significant amount of disk I/O operations, which can result in reduced system performance.
6. **Low Disk Utilization:** Continuous block allocation can lead to low disk utilization, as it may not be possible to allocate smaller portions of the disk to files. This can result in a significant amount of wasted space on the disk.

The linked block allocation scheme is a file allocation scheme in which files are stored on a disk using linked blocks, where each block contains a pointer to the next block in the file. While this scheme has some advantages, it also has several problems, including:

1. **High Overhead:** With linked block allocation, each block in the file requires an additional pointer to the next block, which can result in a significant amount of overhead. This can increase the size of the file and require more disk I/O operations, which can lead to reduced performance.
2. **Poor Performance for Random Access:** With linked block allocation, accessing a specific block in the file requires following a chain of pointers from the beginning of the file to the desired block. This can be slow and inefficient, especially for random access to non-contiguous blocks in the file.
3. **Disk Fragmentation:** With linked block allocation, small files may be scattered across the disk, leading to fragmentation. This can result in wasted space and reduced performance, as the disk needs to seek between non-contiguous blocks to read and write files.
4. **Limited Concurrency:** With linked block allocation, concurrent access to the same file can be difficult, especially if multiple processes are trying to modify the same block in the file. This can result in data inconsistencies and reduced performance.
5. **Security Vulnerabilities:** Linked block allocation can also be vulnerable to security attacks, such as buffer overflow attacks, in which an attacker can modify the pointers in a file to point to malicious code. This can allow the attacker to execute arbitrary code on the system or gain unauthorized access to sensitive data.
6. **Disk Space Overhead:** Linked block allocation also requires additional disk space to store the pointers between blocks, which can reduce the overall disk space utilization and waste disk space.

List any 3 problems of the indexed block allocation schemes

The indexed block allocation scheme is a file allocation scheme in which files are stored on a disk using an index block, which contains pointers to the blocks that make up the file. While this scheme has some advantages, it also has several problems, including:

1. **Limited File Size:** With indexed block allocation, the maximum file size is limited by the size of the index block. If the index block is full, it may not be possible to allocate additional blocks to the file, which can limit the maximum file size.
2. **Disk Fragmentation:** With indexed block allocation, small files may be scattered across the disk, leading to fragmentation. This can result in wasted space and reduced performance, as the disk needs to seek between non-contiguous blocks to read and write files.
3. **Increased Overhead:** With indexed block allocation, each file requires an additional index block to store the pointers to the file's data blocks. This can increase the amount of disk space required to store a file, as well as the amount of disk I/O operations required to read and write the file.

4. **Poor Performance for Random Access:** With indexed block allocation, accessing a specific block in the file requires following a chain of pointers from the index block to the desired block. This can be slow and inefficient, especially for random access to non-contiguous blocks in the file.
5. **Limited Concurrency:** With indexed block allocation, concurrent access to the same file can be difficult, especially if multiple processes are trying to modify the same block in the file. This can result in data inconsistencies and reduced performance.
6. **Disk Space Overhead:** Indexed block allocation also requires additional disk space to store the index block, which can reduce the overall disk space utilization and waste disk space.

**What is a device driver? Write some 7-8 points that correctly describe need, use, placement, particularities of device drivers.**

A device driver is a software program that enables communication between a computer's operating system and a hardware device. The driver acts as a translator, converting commands from the operating system into a language that the device can understand, and translating data from the device into a format that the operating system can use.

The need for a device driver arises because hardware devices have their own unique language and protocols that are specific to their function and manufacturer. The operating system is designed to communicate with devices in a generic way, so it needs a driver to translate the commands and data into the device's specific language.

The use of device drivers allows the operating system to communicate with a wide variety of hardware devices, including printers, scanners, cameras, network adapters, and many others.

Without device drivers, the operating system would not be able to recognize or communicate with these devices, making them useless.

Device drivers are typically placed between the operating system and the hardware device they are designed to control. When the operating system needs to communicate with a device, it sends commands to the driver, which then sends the commands to the device in the correct format. When the device responds with data or status information, the driver translates the information back into a format that the operating system can understand.

Device drivers can be built into the operating system, or they can be separate programs that are loaded when the hardware device is connected or installed. Drivers can be updated or replaced to improve performance or fix bugs, and they are an important part of maintaining a stable and functional computer system.

**What is a zombie process? How is it different from an orphan process?**

-> Zombie processes are the one which are terminated but are waiting for parent to call wait() so that the process will be released from the process table, in case of orphan process parent does not invoke wait() and terminates.

A zombie process is a type of process in a computer system that has completed its execution but still has an entry in the process table. This occurs when the process has sent a message to its parent process to indicate that it has finished executing, but the parent process has not yet read the message to reap its child process. As a result, the process is in a "zombie" state, and its process ID (PID) and other resources are still allocated in the system, but the process itself is not executing.

An orphan process, on the other hand, is a type of process that is still running but has lost its parent process. This can occur when the parent process terminates or crashes before the child process completes its execution. In this case, the child process is re-parented to the init process, which becomes its new parent process. The orphan process continues to run, but it may not be able to access certain resources or perform certain operations without its original parent process.

The main difference between a zombie process and an orphan process is that a zombie process has completed its execution and is waiting for its parent process to reap it, while an orphan process is still running but has lost its parent process. Both types of processes can cause issues in a computer system, as they can consume system resources and cause performance issues. However, zombie processes are generally less problematic than orphan processes, as they are not actively executing and do not consume CPU time or other resources.

### **Which state changes are possible for a process, which changes are not possible?**

->NEW READY RUNNING WAITING TERMINATING. That diagram of states that much possible

A process in a computer system can go through several states during its lifetime. The possible states that a process can be in include:

1. New: The process is being created, and its resources are being allocated.
2. Ready: The process is waiting to be assigned to a processor for execution.
3. Running: The process is being executed by a processor.
4. Waiting: The process is waiting for a certain event or resource, such as in-put/output or a lock.
5. Terminated: The process has finished executing and has been removed from the system.

Some state changes that are possible for a process include:

1. New to Ready: The process has been created and is now waiting to be assigned to a processor for execution.
2. Ready to Running: The process has been assigned to a processor and is now being executed.
3. Running to Waiting: The process is waiting for a certain event or resource.
4. Running to Terminated: The process has finished executing and has been removed from the system.

However, some state changes are not possible or allowed for a process, such as:

1. Running to Ready: A running process cannot go back to the ready state without first waiting for a certain event or resource.
2. Waiting to Running: A waiting process cannot resume execution without first receiving the event or resource it was waiting for.
3. Terminated to any other state: Once a process has finished executing and has been removed from the system, it cannot go back to any other state.

It's important to note that the state changes that are possible for a process may vary depending on the operating system and its implementation.

### **What is mkfs? what does it do? what are different options to mkfs and what do they mean?**

When a storage device is formatted with a file system, the operating system can read and write data to the device, organize the data into files and directories, and track where each file is stored on the device. mkfs creates the necessary structures and metadata for the file system, such as the superblock, inode table, and data blocks.

Some of the different options available with the mkfs command and their meanings are:

1. -t: Specifies the type of file system to create, such as ext2, ext3, ext4, NTFS, FAT32, etc.
2. -c: Checks the device for bad blocks and marks them as unusable.
3. -b: Specifies the block size of the file system. The default value depends on the file system type.
4. -i: Specifies the number of inodes per block group or the inode size.
5. -L: Labels the file system with a name or a volume label.
6. -F: Forces the creation of the file system without prompting for confirmation.
7. -n: Displays the actions that would be taken but does not actually create the file system.

Make file system. mkfs is used to build a Linux file system on a device, usually a hard disk partition. On other operating systems, creating a file system is called formatting.

### **What is the purpose of the PCB? which are the necessary fields in the PCB?**

->Process control block. Has pc,cpu register,cpu scheduling info,memory management info(base+limit), accounting info,i/o info

The PCB (Process Control Block) is a data structure used by operating systems to store information about a running process. The PCB is an essential component of a process management

system, as it provides the operating system with the necessary information to manage and control processes effectively.

The PCB contains a wide range of information about a process, including:

1. Process State: The current state of the process, such as running, ready, waiting, etc.
2. Process ID: A unique identifier assigned to the process by the operating system.
3. Program Counter: The address of the next instruction to be executed.
4. CPU Registers: The values of CPU registers when the process was last interrupted.
5. Memory Management Information: Information about the memory allocated to the process, including its base address, size, and access permissions.
6. I/O Status Information: The status of any open input or output requests made by the process.
7. Accounting Information: Information about the resources used by the process, including CPU time, memory usage, and I/O operations.

The necessary fields in a PCB may vary depending on the operating system and its implementation. However, some of the essential fields that are typically included in a PCB are:

1. Process ID: A unique identifier assigned to the process by the operating system.
2. Process State: The current state of the process, such as running, ready, waiting, etc.
3. Program Counter: The address of the next instruction to be executed.
4. CPU Registers: The values of CPU registers when the process was last interrupted.
5. Memory Management Information: Information about the memory allocated to the process, including its base address, size, and access permissions.
6. I/O Status Information: The status of any open input or output requests made by the process.

The PCB is a critical data structure that allows the operating system to efficiently manage and control processes in a multi-tasking environment.

**Write a C program that uses globals, local variables, static local variables, static global variables, and malloced memory. Compile it. Dump the object code file using `objdump -d -x`. Can you see in the output, the separation into stack, heap, text, etc?**

**Which parts of a C program ((typedef, #define, #include, functions, local vars, statics, globals, #ifdef, ... ) etc occupy RAM when it's a process, and where (code, data, stack, ..)?**

code- main code function code

data-statics, globals,

stack - local vars

heap- mallocated memory

1. #define, #include – no memory

typedef: A typedef declaration does not occupy any memory at runtime, as it is only used by the compiler to define new types.

2. #define: Similarly, a #define statement does not occupy any memory at runtime, as it is only used by the preprocessor to replace text in the code before compilation.

3. #include: An #include statement brings in code from another file, which may contain any of the other components listed here. When the program is compiled, the contents of the included file are merged into the main program's code section.

4. Functions: The code for each function in a C program is stored in the code section of memory, which is typically read-only.

5. Local variables: Local variables are created on the stack when a function is called, and are destroyed when the function returns. They are stored in the stack section of memory.

6. Static variables: Static variables are initialized at program startup and persist throughout the lifetime of the program. They are typically stored in the data section of memory.

7. Global variables: Global variables are also initialized at program startup and persist throughout the lifetime of the program. They are typically stored in the data section of memory.

8. `#ifdef`: `#ifdef` statements are used by the preprocessor to conditionally include or exclude code based on whether a certain macro is defined. The resulting code is merged into the main program's code section.

**Describe the role of CPU (MMU), kernel, compiler in the overall task of "memory management" for a process/program.**

Memory management is a critical aspect of running a process or program, and it involves the cooperation of several components, including the CPU's memory management unit (MMU), the operating system kernel, and the compiler used to build the program.

The role of the CPU's memory management unit (MMU) is to translate virtual memory addresses used by the process into physical memory addresses that correspond to locations in RAM. The MMU does this by using a hardware mechanism called memory mapping, which maps the virtual memory space of the process to a physical address space in RAM. The MMU also provides hardware-based memory protection by preventing a process from accessing memory outside of its allocated virtual address space, which helps to prevent errors and security vulnerabilities.

The kernel of the operating system is responsible for managing the allocation and deallocation of physical memory for all processes running on the system. The kernel maintains a mapping of physical memory pages to virtual memory pages used by each process, and it performs operations such as paging, swapping, and allocating new memory pages as needed. The kernel also provides mechanisms for inter-process communication and synchronization, which may involve sharing memory between processes.

The compiler used to build a program plays a critical role in memory management as well. The compiler determines the layout of the program's code, data, and stack in memory, and it generates machine code that interacts with the MMU and the kernel to allocate and access memory. The compiler may use techniques such as static memory allocation, dynamic memory allocation, and garbage collection to manage the program's memory usage, depending on the programming language and the program's requirements.

In summary, memory management for a process or program is a complex task that requires the cooperation of the CPU's MMU, the operating system kernel, and the compiler used to build the program. The MMU translates virtual memory addresses to physical memory addresses and provides memory protection, the kernel manages the allocation and deallocation of physical memory for all processes, and the compiler determines the layout of the program's memory usage and generates machine code that interacts with the MMU and kernel to allocate and access memory.

**What is the difference between a named pipe and un-named pipe?**

-> unnamed pipe/ ordinary pipe is only used for communication between a child and parent process, while a named pipe there is no need of child and parent communication, follows FIFO. Processes of different ancestry can share data through a named pipe. `bidirectional` . `unnamed` = `unidirectional` `named` = `pipe` act as file

"unnamed" and lasts only as long as the process. A named pipe, however, can last as long as the system is up, beyond the life of the process. It can be deleted if no longer used

A named pipe and an unnamed pipe are two types of interprocess communication (IPC) mechanisms available on Unix-like operating systems. The main difference between these two types of pipes is in how they are created and accessed.

An unnamed pipe, also known as a regular pipe, is a type of IPC that allows data to be transferred between two related processes, typically a parent and a child process. The pipe is created by calling the `pipe` system call, which creates a pipe buffer in memory that is used to transfer data between the processes. An unnamed pipe has no name or file descriptor associated with it and is accessible only to the related processes that share the pipe.

A named pipe, also known as a FIFO (first-in, first-out) or a named FIFO, is a type of IPC that allows data to be transferred between any two or more unrelated processes on the system. A named pipe is created by calling the `mknod` system call, which creates a named file in the file system that



acts as a pipe. The named pipe has a name and is represented by a file descriptor that can be used by any process with appropriate permission to access the pipe.

In summary, the main difference between a named pipe and an unnamed pipe is that a named pipe is accessible to any process on the system that has permission to access it, while an unnamed pipe is accessible only to related processes that share the pipe. A named pipe is created using a file system file, while an unnamed pipe is created using a pipe buffer in memory.

**Describe the steps involved in resolving the path name /a/b/c on an ext2 filesystem.**

1. The path name /a/b/c is broken down into its individual components, separated by the directory separator /.
2. Starting from the root directory /, the ext2 filesystem is searched for the first component a. The root directory is always represented by inode number 2 on an ext2 filesystem.
3. The directory entry for a is located in the root directory, which contains an inode number pointing to the inode of directory a.
4. The inode of directory a is read from disk, and its contents are parsed to find the next component b.
5. The directory entry for b is located in directory a, which contains an inode number pointing to the inode of directory b.
6. The inode of directory b is read from disk, and its contents are parsed to find the final component c.
7. The directory entry for c is located in directory b, which contains an inode number pointing to the inode of file or directory c.
8. The inode of file or directory c is read from disk, and its type and permissions are checked to ensure that the process has the necessary access to the file or directory.
9. If c is a directory, the process may repeat steps 4-8 for any further components in the path name.
10. Once the final component has been resolved, the process can access the file or directory c using its inode number and perform any desired operations.

In summary, resolving a path name on an ext2 filesystem involves recursively traversing the directory tree to locate each component of the path name, until the final file or directory is found and its inode can be used to access its contents.

The maximum possible size of a file on an ext2 partition depends on the block size used by the filesystem. In this case, the block size is 2KB, which means that each file can occupy a maximum number of blocks equal to the total size of the file divided by the block size.

The maximum number of blocks that can be addressed by an ext2 filesystem with a block size of 2KB is  $2^{32} - 1$ , or about 4 billion blocks. This means that the maximum possible size of a file in kilobytes can be calculated as follows:

Maximum file size = (Maximum number of blocks \* Block size) / 1024 = (4 billion \* 2KB) / 1024 = 8,388,608 KB

Therefore, the maximum possible size of a file on an ext2 partition with a block size of 2KB is 8,388,608 KB, or approximately 8.4 GB.

**Write a program that implements a pipe between two processes only (this is Shell assignment-2, first small part).**

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
```

```

int main(void) {
    int fd[2];
    pid_t pid;

    char buf[1024];

    // Create the pipe
    if (pipe(fd) == -1) {
        perror("pipe");
        exit(EXIT_FAILURE);
    }

    // Fork a child process
    pid = fork();

    if (pid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    }

    if (pid == 0) { // Child process
        close(fd[1]); // Close the write end of the pipe

        while (read(fd[0], buf, sizeof(buf)) > 0) { // Read from the read end of the pipe
            printf("Child received message: %s", buf);
        }

        close(fd[0]); // Close the read end of the pipe
        exit(EXIT_SUCCESS);
    } else { // Parent process
        close(fd[0]); // Close the read end of the pipe

        char *msg = "Hello, child process!\n";
        write(fd[1], msg, strlen(msg)); // Write to the write end of the pipe

        close(fd[1]); // Close the write end of the pipe
        exit(EXIT_SUCCESS);
    }
}

```

### **What are privileged instructions? Why are they required? What is their relationship to 2 modes of CPU execution?**

2 modes kernel mode and user mode . Priviledged instruction only run by kernel.kernel=0 user mode bit=1.protection coz of dualmode

Privileged instructions are CPU instructions that can only be executed in privileged mode. These instructions are designed to perform operations that are critical to the functioning of the operating system, such as manipulating hardware registers, controlling system re-sources, and accessing privileged data structures. Examples of privileged instructions in-clude loading or modifying control registers, executing I/O operations, and changing the memory protection level.

Privileged instructions are required to ensure the safety and stability of the operating sys-tem. By restricting access to certain CPU instructions and system resources, the operating system can prevent user-level processes from interfering with critical system functions and compromising the

security of the system. For example, if a user-level process were able to execute privileged instructions directly, it could potentially cause the system to crash or allow malicious actors to gain control of the system.

Most CPUs support two modes of execution: user mode and privileged mode (also known as kernel mode). User mode is the mode in which most applications and processes run, and it provides restricted access to system resources and privileged instructions. Privileged mode, on the other hand, is the mode in which the operating system kernel runs, and it provides full access to system resources and privileged instructions.

In privileged mode, the CPU can execute privileged instructions and access system re-sources that are restricted in user mode. By contrast, in user mode, the CPU is limited to executing a subset of instructions that do not require privileged access. This separation of privileges between user mode and privileged mode helps to ensure the security and sta-bility of the system by preventing user-level processes from directly accessing sensitive system functions and resources.

### **Explain what happens in pre-processing, compilation and linking and loading.**

In general, the process of compiling and linking a program involves several stages: pre-processing, compilation, linking, and loading. Here's a brief explanation of what happens at each stage:

1. **Pre-processing:** During this stage, the pre-processor takes the source code of a program and performs various operations on it, such as including header files, re-placing macros with their definitions, and removing comments. The output of this stage is a modified version of the original source code, which is then passed on to the compiler.
  2. **Compilation:** During the compilation stage, the compiler takes the pre-processed source code and translates it into machine code, which is the binary code that can be executed by the CPU. This involves several sub-stages, including lexical analysis (breaking the code into tokens), syntax analysis (parsing the tokens into a syntax tree), semantic analysis (checking the syntax tree for correctness), code generation (generating machine code from the syntax tree), and optimization (improving the efficiency of the generated code). The output of this stage is an object file, which contains the machine code for a single compilation unit.
  3. **Linking:** During the linking stage, the linker takes one or more object files and combines them into a single executable file or shared library. This involves several sub-stages, including symbol resolution (matching up function and variable refer-ences across different object files), relocation (adjusting the machine code of the object files to work together), and library linking (linking in external libraries that are required by the program). The output of this stage is an executable file or shared library, which can be executed or loaded by the operating system.
  4. **Loading:** During the loading stage, the operating system loads the executable file or shared library into memory and prepares it for execution. This involves several sub-stages, including memory allocation (allocating memory for the program and its data), symbol binding (mapping symbols to their memory addresses), and relo-cation (adjusting the machine code to work with the actual memory layout). Once the loading is complete, the program can be executed by the CPU.
- Overall, the process of compiling and linking a program involves several complex stages, each of which is responsible for a different aspect of the process. By breaking the process down into these stages, compilers and linkers can perform each stage more efficiently and accurately, leading to faster and more reliable program builds.

### **Why is the kernel called an event-drive program?**

The kernel is often referred to as an event-driven program because it is designed to respond to events that occur in the system, such as hardware interrupts, software interrupts, and system calls. These events trigger specific actions or routines within the kernel, allowing it to manage system resources and provide services to user-level processes.

For example, when a process makes a system call, such as opening a file or creating a new process, it triggers an event in the kernel. The kernel then responds by executing the appropriate system call routine, which performs the requested operation and returns the result to the calling process.

Similarly, when a hardware interrupt occurs, such as a keyboard input or disk I/O completion, it triggers an event in the kernel. The kernel then responds by executing the appropriate interrupt handler routine, which performs the necessary operations to handle the interrupt and resume normal system operation.

In an event-driven program like the kernel, the flow of control is driven by events rather than by a sequential program flow. This allows the kernel to be more responsive and efficient, since it only needs to execute the code required to handle the specific events that occur in the system.

### **What are the limitations of segmentation memory management scheme?**

Segmentation is a memory management scheme that divides a process's logical address space into segments, each of which represents a portion of the process's memory. While segmentation has some advantages over other memory management schemes, it also has several limitations, including:

**External fragmentation:** Segmentation can lead to external fragmentation, where the memory space becomes divided into small, non-contiguous blocks, making it difficult to allocate new memory segments of a particular size.

**Overhead:** Segmentation requires additional overhead in terms of maintaining information about each segment, such as its size, location, and access permissions. This can lead to increased memory usage and slower performance.

**Security:** Segmentation can make it more difficult to enforce security policies, since segments may overlap or have different access permissions, making it harder to control access to specific portions of memory.

**Limited scalability:** Segmentation may not scale well for very large address spaces or very large numbers of processes, since the overhead required to manage the segments can become prohibitively large.

**Complexity:** Segmentation can add complexity to the memory management system, making it more difficult to implement and debug.

**Difficulty of sharing memory:** Since segments are not necessarily contiguous in physical memory, sharing memory between processes can be more difficult than with other memory management schemes, such as paging.

Overall, while segmentation can be a useful memory management scheme in certain situations, it also has several limitations that must be taken into account when designing a memory management system.

### **How is the problem of external fragmentation solved?**

External fragmentation is a common problem in memory management schemes where the memory space becomes divided into small, non-contiguous blocks, making it difficult to allocate new memory segments of a particular size. There are several ways to address this problem:

**Compaction:** Compaction is a technique where the operating system periodically moves allocated segments of memory so that all free memory becomes contiguous. This can reduce external fragmentation, but it can also be time-consuming and may require the relocation of running processes, which can impact system performance.

**Buddy allocation:** Buddy allocation is a memory allocation technique that divides memory into fixed-sized blocks, and groups them into pairs that are contiguous in memory. When a request for memory is made, the system allocates the smallest block size that will meet the request, and splits a larger block if necessary. This can reduce external fragmentation by ensuring that free blocks are always of a specific size.

**Memory paging:** Paging is a memory management scheme that divides the memory space into fixed-sized pages, and maps logical addresses to physical addresses using a page table. Paging can reduce external fragmentation by allowing the operating system to allocate memory on demand, and by allowing the system to move pages around in memory to create contiguous blocks.

Virtual memory: Virtual memory is a memory management scheme that allows processes to use more memory than is physically available in the system. This is achieved by using a combination of paging and demand paging, where only the pages that are actually needed are loaded into physical memory. Virtual memory can help reduce external fragmentation by allowing the system to allocate memory more efficiently, and by allowing processes to use more memory without causing external fragmentation.

### **Does paging suffer from fragmentation of any kind?**

Paging can suffer from internal fragmentation, which is different from external fragmentation that can occur in other memory management schemes such as segmentation.

Internal fragmentation occurs when the size of the page allocated to a process is larger than the size of the memory segment that the process needs. In this case, some portion of the page is left unused, resulting in inefficient use of memory. This is because paging divides the memory into fixed-size pages, and a process is allocated one or more whole pages, even if the process does not require all the space in the page.

To reduce internal fragmentation, operating systems typically use demand paging, where pages are only loaded into physical memory when they are actually needed. This means that only the required portion of the page is used, reducing internal fragmentation. Additionally, some systems use a technique called page coloring to allocate pages in such a way that minimizes internal fragmentation.

Overall, while paging can suffer from internal fragmentation, this problem is typically less severe than the external fragmentation that can occur in other memory management schemes. Additionally, demand paging and other techniques can be used to minimize internal fragmentation and make paging a more efficient memory management scheme

### **in this program: `int main() { int a[16], i; for(i = 0; ; i++) printf("%d\n", a[i]); }` why does the program not segfault at `a[16]` or some more values?**

The program does not necessarily segfault at `a[16]` or some more values because the behavior of the program is undefined.

The program declares an integer array "a" of size 16, but it does not initialize the elements of the array. Therefore, the initial values of the elements of the array are indeterminate, and the behavior of accessing them is undefined.

In practice, when the program is executed, the elements of the array may contain random values that happen to be valid integers, and the program may continue to print the values until it encounters an invalid memory location or some other error. Alternatively, the program may crash or produce unpredictable results when it tries to access invalid memory locations.

In general, accessing memory beyond the bounds of an array is a serious programming error and can lead to unexpected behavior and security vulnerabilities. To avoid such errors, it is important to always initialize arrays before using them and to ensure that all array accesses are within the bounds of the array.

### **What is segmentation fault? why is it caused? who detects it and when? how is the process killed after segmentation fault?**

A segmentation fault (often abbreviated as "segfault") is an error that occurs when a program attempts to access memory that it is not allowed to access, typically because the memory address does not exist or is outside the program's allocated memory space.

A segmentation fault can be caused by a variety of programming errors, such as dereferencing a null pointer, accessing an array out of bounds, or attempting to execute code in a non-executable memory region. In general, any attempt to access memory that is not allocated or outside the program's allocated space can result in a segmentation fault.

The operating system is responsible for detecting segmentation faults. When a program attempts to access invalid memory, the CPU generates a fault or exception, which is detected by the operating

system's memory management unit. The operating system then sends a signal to the program indicating that a segmentation fault has occurred.

By default, when a program receives a segmentation fault signal, it is terminated immediately by the operating system. The process may generate a core dump, which is a file containing a snapshot of the program's memory at the time of the segmentation fault. This core dump can be analyzed by developers to diagnose the cause of the segmentation fault and fix the underlying programming error.

In some cases, programs can catch and handle segmentation faults themselves, typically by using signal handlers. However, this is generally not recommended, as segmentation faults indicate serious programming errors that should be fixed rather than ignored.

### **What is the meaning of "core dumped"?**

#### **Core dump=memory related**

"Core dumped" is a message that may appear in the terminal or console after a program crashes or is terminated due to a segmentation fault or other fatal error.

When a program crashes, the operating system may create a file called a "core dump" or "core file", which contains a snapshot of the program's memory at the time of the crash. This core dump file can be useful for developers and system administrators to diagnose the cause of the crash and fix any bugs or errors in the program.

The message "core dumped" typically indicates that a core dump file has been created. The file may be named "core", "core.pid", or similar, where "pid" is the process ID of the crashed program. The location of the core dump file may vary depending on the operating system and system configuration.

In general, when a program crashes and generates a core dump file, it is important to analyze the file to identify the cause of the crash and fix any errors or bugs in the program. This can help to prevent similar crashes and improve the reliability and stability of the program.

### **What is voluntary context switch and non-voluntary context switch on Linux? Name 2 processes which have lot of non-voluntary context switches compared to voluntary.**

In Linux, a context switch refers to the process of saving the current state of a running process (i.e. its registers and other CPU state) and restoring the state of another process so that it can run.

A voluntary context switch occurs when a process voluntarily yields the CPU, typically by calling a system call like `sleep()` or `yield()`. In this case, the process explicitly requests to be put to sleep or to allow other processes to run, so the context switch is considered voluntary.

A non-voluntary context switch occurs when a running process is preempted by the operating system or another process, typically due to a higher-priority process becoming runnable or due to a hardware interrupt. In this case, the context switch is considered non-voluntary because the running process did not explicitly request to be switched out.

Some processes that may have a lot of non-voluntary context switches compared to voluntary context switches include:

Interactive programs that perform I/O operations, such as web servers, database servers, or file servers. These programs may spend a lot of time waiting for I/O operations to complete, which can result in frequent non-voluntary context switches.

Real-time programs that have strict timing requirements, such as audio or video processing applications. These programs may be preempted by the operating system or other processes if they do not complete their tasks within a certain timeframe, resulting in frequent non-voluntary context switches.

It's worth noting that the ratio of voluntary to non-voluntary context switches can vary widely depending on the nature of the workload and the system configuration.

**why does the Wikipedia entry for ext2 say this: "Max. file size: 16 GiB – 2 TiB" ?**

The Wikipedia entry for ext2 mentions a maximum file size of "16 GiB – 2 TiB" because the maximum file size that can be supported by an ext2 file system depends on several factors, including the block size and the specific implementation of the file system.

In general, the maximum file size on an ext2 file system is determined by the maximum number of blocks that can be addressed by the file system. Each block is typically 1KB, 2KB, or 4KB in size, and the maximum number of blocks that can be addressed depends on the file system's addressing scheme.

For ext2 file systems with a block size of 1KB, the maximum file size is 16GB. For file systems with a block size of 2KB, the maximum file size is 256GB. For file systems with a block size of 4KB, the maximum file size is 2TB.

However, it's worth noting that these values are theoretical limits and may not be achievable in practice due to other factors such as disk space limitations, file system overhead, and system limitations. Additionally, newer versions of the ext file system, such as ext3 and ext4, may support larger maximum file sizes.